

# Deep learning

CNN and supervised learning

March 2019 @ TELECOM Nancy  
T. BAGREL

# Introduction

Interesting NN architectures (supervised)

## Convolutional Neural Network (CNN)

- ▶ Supervised :(
- ▶ Most common architecture for image processing
- ▶ Position-invariant feature extraction (with shared weights)
- ▶ Very efficient

# Introduction

Interesting NN architectures (unsupervised)

## Restricted **B**oltzmann **M**achine (RBM)

- ▶ Unsupervised!
- ▶ General feature extraction
- ▶ Basic block of **D**eep **B**elief **N**etworks

## Autoencoders (AE)

- ▶ Unsupervised!
- ▶ Can be used for denoising
- ▶ Can be used for general feature extraction

# Autoencoders

## Structure and principle

### Principle

- ▶ **Simple** : reproduce the input at the output
- ▶ **How** : pass the input (image) through the NN, compute error from  $\text{diff}(I, O)$ , adjust weights with GD

### Feature extraction

**Bottleneck** : “center” of the NN where there is the lowest number of neurons. Activation in the “bottleneck” for a given input (image) represents its compressed form.

# Autoencoders

## Denoising




FIGURE – Autoencoder with feature extraction purpose

## Denoising

Denoising can be achieved with roughly the same architecture

# CNN Components

## Gradient Descent

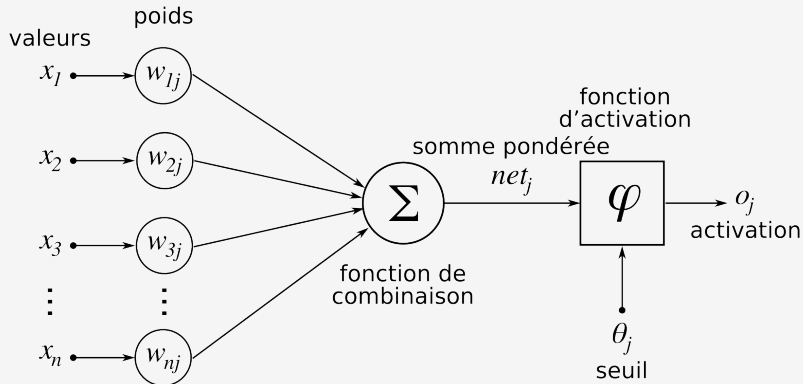


FIGURE – Neuron (perceptron) schema

# CNN Components

## Gradient Descent

### Problem solved by GD

In which direction do we need to update our weights to minimize the error?

### Last layer neuron $n$ situation

$E$  : error,  $L$  : loss function,  $e$  : expected value,  $o$  : obtained value,  $\varphi$  : activation function,  $x_{n,\cdot}$  : inputs for the last layer neuron,  $w_{n,\cdot}$  : weights for the last layer neuron

$$E = L(e, o) \tag{1}$$

$$= L \left( e, \varphi \left( \sum_k x_{n,k} w_{n,k} \right) \right) \tag{2}$$

# CNN Components

## Gradient Descent

### What we need to know

- ▶ **loss function**  $L(e, o)$  and  $\frac{\partial L}{\partial o}$
- ▶ **activation function**  $\varphi(r)$  and  $\frac{d\varphi}{dr}$

### What we can compute

With the **chain rule**, we can compute  $\frac{\partial E}{\partial w_{n,k}}(e, o) \quad (\forall k)$




# CNN Components

## Gradient descent and adjustment made

Weights adjustment ( $\alpha$  : learning rate)

$$w_{n,k} \leftarrow w_{n,k} - \alpha \cdot \frac{\partial E}{\partial w_{n,k}}(e, o)$$



# CNN Components

## Learning rate and advanced adjustment methods




FIGURE – Importance of an appropriate learning rate

# CNN Components

## Advanced adjustment methods

Called optimizers in **TensorFlow**

### Roles

- ▶ Try to prevent staying in a non-global local minima during learning
- ▶ **Examples** : Nesterov Momentum, Adam...

### Used in my NNs

I use Nesterov's momentum :  
`optimizers.SGD(nesterov=True, momentum=M)`  
because it's recommended on most tutorials

# CNN Components

## GD vs SGD vs BGD

Each one is a different manner to apply Gradient Descent

### Differences


- ▶ **(basic) Gradient Descent** : weights are updated after every epoch (whole dataset pass) → **very costly** but more stable than SGD!
- ▶ **Stochastic Gradient Descent** : weights are updated after **every** sample → **very fast**!
- ▶ **Batch Gradient Descent** : weights are updated after N samples ( $10 \leq N \leq 100$  : batch size) → **best of both worlds**!

# CNN Components

## Convolutional layer

The (3D) **dot product** between the **filter weights** and an **input image zone** of the same size/shape gives an activation value. Doing this operation for all possible image zones (with **overlap**) gives the **activation map** for this filter.

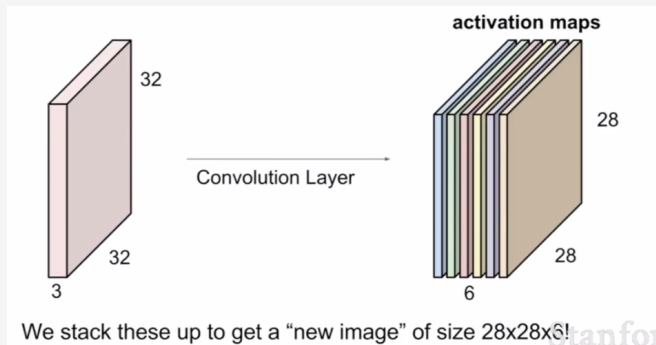
*It explains the position invariance of feature extraction for CNNs!*



# CNN Components

## Convolutional layer


A **convolutional layer** is composed of several filters, each one with its **own weights**. Hence, the output of the convolutional layer is a **stack of activation maps** (one for each filter)



# CNN Components

## Pooling layer

A **pooling layer** is used to downsize/down-sample its input. **All** the values in a zone of the input volume results in **one** value in the output volume, by application of a simple function. Input zones **must not overlap** here!  
A very common example is the *max pooling*  $2 \times 2$  :



# CNN Components

## Output layer for classification

The output layer is used to extract classification information from a final dense/convolutional/pooling layer.

### Types

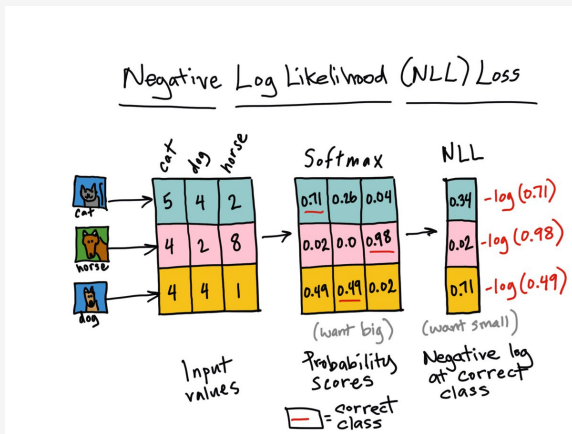
- ▶ For binary classification : Dense layer with **sigmoid** activation function and **1** neuron (1 output value)
- ▶ For multiclass classification : Dense layer with **softmax** activation function and  **$N_{\text{classes}}$**  neurons, followed by a  $\text{arg\_max}(\cdot \cdot \cdot)$  function



# CNN Components

## Loss function for classification

**Cross entropy** (aka. Negative Log Likelihood) is the loss function used in classification, which works well with the **sigmoid** or **softmax** activation functions (output layer)



# CNN Components

## CNN schema for multiclass




FIGURE – CNN architecture

# Medical imaging

## Availability

### The issue

There are medical image datasets available online, but

- ▶ High resolution
- ▶ Difficult to find the differences or the good class for an untrained eye
- ▶ File format and dataset organisation is not always normalized

### Examples

- ▶ List of public datasets : [bit.ly/2E0ziVp](https://bit.ly/2E0ziVp), [bit.ly/2u0Da0Q](https://bit.ly/2u0Da0Q)
- ▶ Datasets : [bit.ly/2F34jqb](https://bit.ly/2F34jqb) (cancer cells), [bit.ly/2u99w9L](https://bit.ly/2u99w9L) (neuroimaging), [bit.ly/2cGnNQL](https://bit.ly/2cGnNQL) (Alzheimer)...

# Medical imaging


What to do so?

## Requirements

- ▶ Large dataset
- ▶ Low-res images (max. 256 x 256)
- ▶ Something which looks like medical images

## Idea

Generate **brain 2D images** and add random “tumors” on them



# brain\_tumor\_factory

Let's get our hands dirty!

## How to generate?

Home-made C program : `brain_tumor_factory`

- ▶ take a `single` base image
- ▶ add tumor(s) randomly depending on the chosen scenario
- ▶ alter the resulting image, to ensure that each output image is a bit unique

Program size :  $\simeq$  500 LOC

Output :  $\simeq$  1000 images/s


Image size : `96 x 128 x 1` (grayscale)

# brain\_tumor\_factory

Creating tumors : tumor function

## Instructions

Take a circle of radius  $R_{\text{area}}$ , and put  $n \in \llbracket n_{\min}, n_{\max} \rrbracket$  “soft discs” of radius  $r_i \in [r_{\min}, r_{\max}]$  with their center placed randomly inside the chosen area



# brain\_tumor\_factory

Altering the resulting image : `alter_full` function

## Alterations

- ▶ `shrink` : shrink the image on x-axis and/or y-axis
- ▶ `move` : shift the image on x-axis and/or y-axis
- ▶ (`ghost_blur`  $0.1\%$ ) : simulate motion blur on a shot
- ▶ `alter_contrast` :  $\pm$ [brightness range], with custom "brightness center"
- ▶ `alter_shade` : global flat  $\pm$ [brightness]
- ▶ `alter_rdc` : per-pixel flat random  $\pm$ [brightness]

For each image, alterations parameters are random numbers taken from a `uniform` or `normal` distribution

# brain\_tumor\_factory


Base image





# brain\_tumor\_factory

Altered images




# brain\_tumor\_factory

## Syndrome A

### Description

- ▶  $\llbracket 1, 2 \rrbracket$  big tumor(s)
- ▶  $\llbracket 0, 2 \rrbracket$  medium tumor(s)




# brain\_tumor\_factory

## Syndrome B

### Description

- ▶  $\llbracket 4, 8 \rrbracket$  medium tumor(s)
- ▶  $\llbracket 0, 3 \rrbracket$  small/tiny tumor(s)




# brain\_tumor\_factory

## Syndrome C


### Description

- ▶ only  $\llbracket 20, 48 \rrbracket$  small/tiny tumor(s)



# brain\_tumor\_factory

Guess game : EASY mode



# brain\_tumor\_factory

Guess game : HARD mode

## Observation

It turns out that it's kinda easy to find the correct syndrome type, isn't it?

Let's try the **HARD** mode then (more realistic) 😈


## HARD mode specifications

- ▶ Tumor opacity **greatly** reduced
- ▶ Tumor size reduced a bit

The **base image** will always be placed on the **LEFT** of the image whose type must be determined


# brain\_tumor\_factory

Guess game : HARD mode



# brain\_tumor\_factory

Guess game : HARD mode





Answer

Syndrome C




# brain\_tumor\_factory

Guess game : HARD mode



# brain\_tumor\_factory

Guess game : HARD mode




Answer

Syndrome A


# brain\_tumor\_factory

Guess game : HARD mode



# brain\_tumor\_factory

Guess game : HARD mode




Answer

Z (healthy patient)


# brain\_tumor\_factory

Guess game : HARD mode



# brain\_tumor\_factory

Guess game : HARD mode




Answer

Syndrome B


# brain\_tumor\_factory

Guess game : HARD mode



# brain\_tumor\_factory

Guess game : HARD mode





Answer

Z (healthy patient)




# brain\_tumor\_factory

Guess game : HARD mode



# brain\_tumor\_factory

Guess game : HARD mode



Answer

Syndrome C

# brain\_tumor\_factory

Guess game : HARD mode

Observation

Not so easy this time? 😈

Let's see what **deep learning** can do

# brain\_tumor\_nn

## Description

### Framework

- ▶ DL4j / ScalNet (Scala)
- ▶ not really fast (afterwards)
- ▶ Keras-like API

### Architecture

- ▶ 3 Conv2D (MaxPooling2D inbetween) : 16 64 256
- ▶ Dense 1024
- ▶ GD (BGD) with Nesterov Momentum
- ▶ (adapted from MNIST NN examples)

# brain\_tumor\_nn

## Description

### Issue (after 2 hours learning)

- ▶ Accuracy blocked at 0.25
- ▶ loss wasn't decreasing
- ▶ a **single type** predicted everytime

### Why (afterwards)

- ▶ I didn't wait long enough
- ▶ 3 Conv layer NN + slow framework didn't help much neither
- ▶ in addition, learning on mobile CPU (i7) is still slow
- ▶ (but afterwards) architecture wasn't really the issue

